

Take your microservices to the next level with gRPC

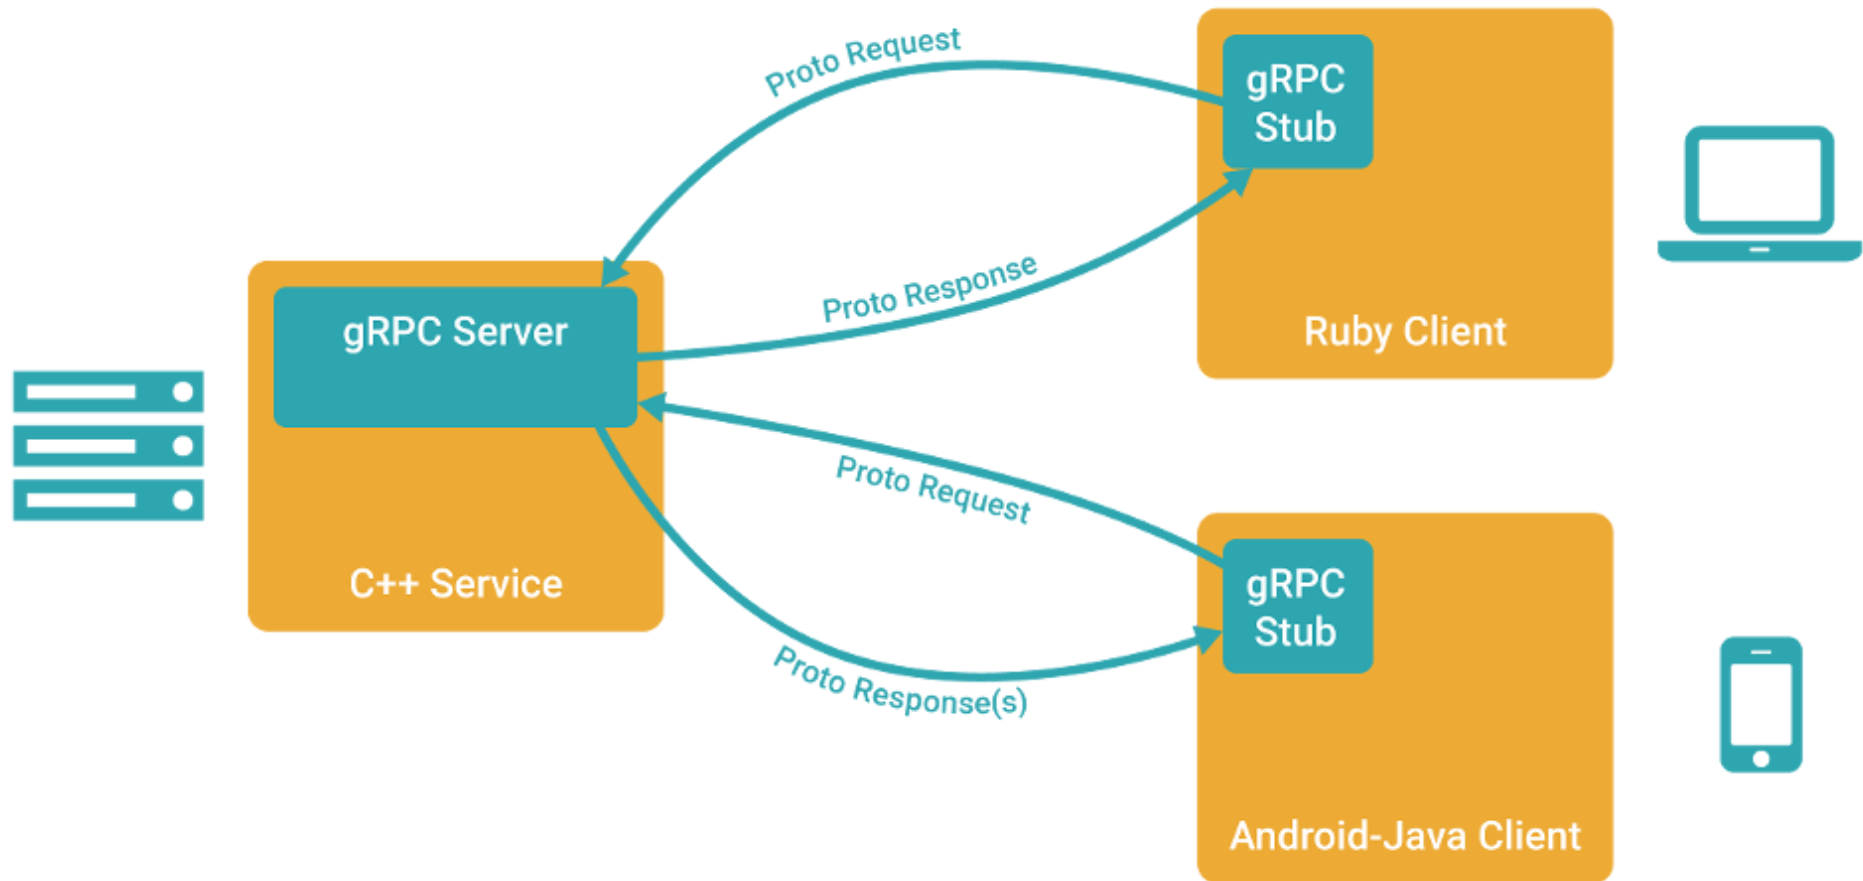
Mateusz Dymiński

gRPC nowy framework RPC stworzony przez Google

grpc.io (<http://grpc.io>)

- RPC oraz strumieniowane RPC
- Wspierane języki: **Java, Go, C, C++, Node.js, Python, Ruby, Objective-C, PHP, i C#**
- IDL: **Protocol Buffers 3**
- Transport: **HTTP2**
- Auth: **SSL/TLS**

gRPC



Strumieniowanie RPC w obu kierunkach

Klient rozpoczyna komunikację z serwerem.

Wiadomości dostarczane są w kolejności FIFO.

Wiele równoległych strumieni.

Jedno połączenie - dzięki HTTP/2

Zapewnia buforowanie oraz kontrolę przepływu.

Przykłady:

- strumieniowanie w obu kierunkach: chat
- strumień serwer → klient: wykres giełdowy
- strumień klient → serwer: agregowanie danych z sensora

gRPC użytkownicy

150+ importów - [google.golang.org/grpc](https://godoc.org/google.golang.org/grpc) (<https://godoc.org/google.golang.org/grpc?importers>) on godoc.org (<http://godoc.org>)

Biblioteki w Go:

- [CockroachDB](https://github.com/cockroachdb/cockroach) (<https://github.com/cockroachdb/cockroach>): Super stable distributed DB
- [Bazil](http://bazil.org) (<http://bazil.org>): distributed file system
- [CoreOS/Etcd](http://coreos.com/etcd/) (<http://coreos.com/etcd/>): distributed consistent key-value store
- [Google Cloud Bigtable](https://godoc.org/google.golang.org/cloud/bigtable) (<https://godoc.org/google.golang.org/cloud/bigtable>): sparse table storage
- [YouTube/Vitess](http://vitess.io/) (<http://vitess.io/>): storage platform for scaling MySQL
- [gRPC Gateway](https://github.com/gengo/grpc-gateway) (<https://github.com/gengo/grpc-gateway>): revers proxy - translates REST into gRPC

Mikroserwisy - problemy

- JSON
- Wersjonowanie API
- Zmiany w modelach w obrębie wielu serwisów
- Śledzenie wywołań
- QoS - deadline
- Anulowanie wywołań

Mikroserwisy - problemy - JSON

MacBook Pro 2.6 GHz i7 16GB

protobuf vs golang-json-serializer

test	iter	time/iter	bytes alloc	allocs

BenchmarkJsonMarshal-8	500000	3714 ns/op	1232 B/op	10 allocs/op
BenchmarkJsonUnmarshal-8	500000	4125 ns/op	416 B/op	7 allocs/op
BenchmarkProtobufMarshal-8	1000000	1554 ns/op	200 B/op	7 allocs/op
BenchmarkProtobufUnmarshal-8	1000000	1055 ns/op	192 B/op	10 allocs/op

encoded sizes:

default 1232

protobuf 243

Diff serialization: $3714 / 1554 = 2.39$

Diff deserialization: $4125 / 1055 = 3.91$

Mikroserwisy - problemy - JSON

Macbook Pro 2.7 GHz i7 16GB

protobuf vs jackson

test	min	max	avg

jackson serialization	53.3	78.7	62.5
jackson deserialization	110.1	130.9	114.9
protobuf serialization	10.2	82.1	19.0
protobuf deserialization	19.3	35.2	25.1
encoded sizes:			
jackson	949		
protobuf	258		

Diff serialization: $62.5 / 19.0 = 3.29$

Diff deserialization: $114.9 / 25.1 = 4.58$

Mikroserwisy - problemy

<http://some-host.com/users/mdyminski> - **GET? POST? PUT? DELETE?**

VS


`deleteUser(String name)`

Mikroserwisy - problemy JSON

- Wolny
- Brak schematu
- Rozmiar
- Wersjonowanie
- Walidacja

Dummy Google

Architektura



Demo! - Google search

Protocol definition

```
syntax = "proto3";
option java_multiple_files = true;
option java_package = "com.grpc.search";
option java_outer_classname = "SearchProto";
option objc_class_prefix = "GGL";
package search;

service Google {
    // Search returns a Google search result for the query.
    rpc Search(Request) returns (Result) {}

    // Watch returns a stream of Google search results for the query.
    rpc Watch(Request) returns (stream Result) {}
}

message Request {
    string query = 1;
}

message Result {
    string title = 1;
    string url = 2;
    string snippet = 3;
}
```

Wygenerowany kod - Golang

```
type GoogleClient interface {  
    // Search returns a Google search result for the query.  
    Search(ctx context.Context, in *Request, opts ...grpc.CallOption) (*Result, error)  
    // Watch returns a stream of Google search results for the query.  
    Watch(ctx context.Context, in *Request, opts ...grpc.CallOption) (Google_WatchClient, error)  
}
```

```
type GoogleServer interface {  
    // Search returns a Google search result for the query.  
    Search(context.Context, *Request) (*Result, error)  
    // Watch returns a stream of Google search results for the query.  
    Watch(*Request, Google_WatchServer) error  
}
```

```
type Request struct {  
    Query string `protobuf:"bytes,1,opt,name=query" json:"query,omitempty"`  
}
```

```
type Result struct {  
    Title   string `protobuf:"bytes,1,opt,name=title" json:"title,omitempty"`  
    Url     string `protobuf:"bytes,2,opt,name=url" json:"url,omitempty"`  
    Snippet string `protobuf:"bytes,3,opt,name=snippet" json:"snippet,omitempty"`  
}
```

Wygenerowany kod - Java

```
public static final int TITLE_FIELD_NUMBER = 1;
private volatile java.lang.Object title_;
/**
 * <code>optional string title = 1;</code>
 */
public java.lang.String getTitle() {
    java.lang.Object ref = title_;
    if (ref instanceof java.lang.String) {
        return (java.lang.String) ref;
    } else {
        com.google.protobuf.ByteString bs =
            (com.google.protobuf.ByteString) ref;
        java.lang.String s = bs.toStringUtf8();
        title_ = s;
        return s;
    }
}
```



Wygenerowany kod - Java

```
/**
 * Creates a new async stub that supports all call types for the service
 */
public static GoogleStub newStub(io.grpc.Channel channel) {
    return new GoogleStub(channel);
}

/**
 * Creates a new blocking-style stub that supports unary and streaming output calls on the service
 */
public static GoogleBlockingStub newBlockingStub(
    io.grpc.Channel channel) {
    return new GoogleBlockingStub(channel);
}

/**
 * Creates a new ListenableFuture-style stub that supports unary and streaming output calls on the service
 */
public static GoogleFutureStub newFutureStub(
    io.grpc.Channel channel) {
    return new GoogleFutureStub(channel);
}
```


Frontend odpala Search na serwerach i czeka na pierwszy wynik



Demo klient

- Frontend śledzenie wywołań
- Backend śledzenie wywołań
- Logi połączenia

Kod źródłowy - klient



Kod źródłowy - klient (konstruktor)

```
private final ManagedChannel channel;
private final GoogleGrpc.GoogleBlockingStub googleBlockingStub;

/**
 * Construct client with blocking API to Frontend server at {@code host:port}.
 */
public Client(String host, int port) {
    channel = ManagedChannelBuilder.forAddress(host, port)
        .usePlaintext(true)
        .build();
    googleBlockingStub = GoogleGrpc.newBlockingStub(channel);
}
```

Kod źródłowy - klient (search)

```
/**
 * Search query in dummy google backend.
 */
public Result search(String query) {
    logger.info("Starting search for " + query + " ...");

    final Request request = Request.newBuilder().setQuery(query).build();

    return Try
        .ofFailable(() -> googleBlockingStub.search(request))
        .onSuccess(r -> logger.info("Search result: " + r))
        .onFailure(e -> logger.log(Level.SEVERE, "RPC failed: {0}", e))
        .getUnchecked();
}
```

gRPC blokuje wątek w momencie wywołania.

gRPC propaguje anulowanie żądania po otrzymaniu pierwszego rekordu.


Kod źródłowy - klient nieblokujący (search)

```
/**
 * Search query in async way in dummy google backend.
 */
public Result search(String query) {
    logger.info("Starting search for " + query + " ...");

    Request request = Request.newBuilder().setQuery(query).build();
    ListenableFuture<Result> resultFuture = googleFutureClient.search(request);

    return Try.ofFailable(() -> resultFuture.get(1000, TimeUnit.MILLISECONDS))
        .onSuccess(r -> logger.info("Search result: " + r))
        .onFailure(e -> logger.log(Level.SEVERE, "RPC failed: {0}", e))
        .getUnchecked();
}
```

Kod źródłowy - backend



Kod źródłowy - backend

```
public class Backend {  
  
    .  
    .  
    .  
  
    /**  
     * Main launches the server from the command line.  
     */  
    public static void main(String[] args) throws IOException, InterruptedException {  
        String id = "0";  
  
        if (args.length > 0) {  
            id = args[0];  
        }  
  
        final Backend server = new Backend();  
        server.start(Integer.parseInt(id));  
        server.blockUntilShutdown();  
    }  
}
```

Kod źródłowy - backend

```
private int port = 36061;
private Server server;

private void start(int id) throws IOException {
    server = ServerBuilder.forPort(port)
        .addService(new GoogleImpl(id))
        .build().start();
    logger.info("Server started, listening on " + port);
    Runtime.getRuntime().addShutdownHook(new Thread() {
        @Override
        public void run() {
            Backend.this.stop();
        }
    });
}

private void stop() {
    if (server != null) { server.shutdown(); }
}

private void blockUntilShutdown() throws InterruptedException {
    if (server != null) { server.awaitTermination(); }
}
```

Kod źródłowy - backend


```
private class GoogleImpl extends GoogleGrpc.AbstractGoogle {
    private final int id;

    public GoogleImpl(int id) {
        this.id = id;
    }

    @Override
    public void search(Request req, StreamObserver<Result> responseObserver) {
        sleepRandTime();
        Result reply = Result
            .newBuilder()
            .setTitle(format("result for [%s] from backend %d", req.getQuery(), id))
            .build();
        responseObserver.onNext(reply);
        responseObserver.onCompleted();
    }

    private void sleepRandTime() {
        Try.ofFailable(() -> {
            Thread.sleep((new Random().nextInt(9) + 1) * 10;
        }).getUnchecked();
    }
}
```

Kod źródłowy - frontend



Kod źródłowy - frontend (search)

Search zwraca wynik tak szybko jak dostanie pierwszy wynik z wywołań. gRPC anuluje pozostałe wywołania backend. Search RPC za pomocą ctx:

```
func (s *server) Search(ctx context.Context, req *pb.Request) (*pb.Result, error) {
    c := make(chan result, len(s.backends))
    for _, b := range s.backends {
        go func(backend pb.GoogleClient) {
            res, err := backend.Search(ctx, req)
            c <- result{res, err}
        }(b)
    }
    first := <-c
    return first.res, first.err
}
```

```
type result struct {
    res *pb.Result
    err error
}
```

Strumieniowanie RPC

Nowa metoda Watch

```
syntax = "proto3";
option java_multiple_files = true;
option java_package = "com.grpc.search";
option java_outer_classname = "SearchProto";
option objc_class_prefix = "GGL";
package search;

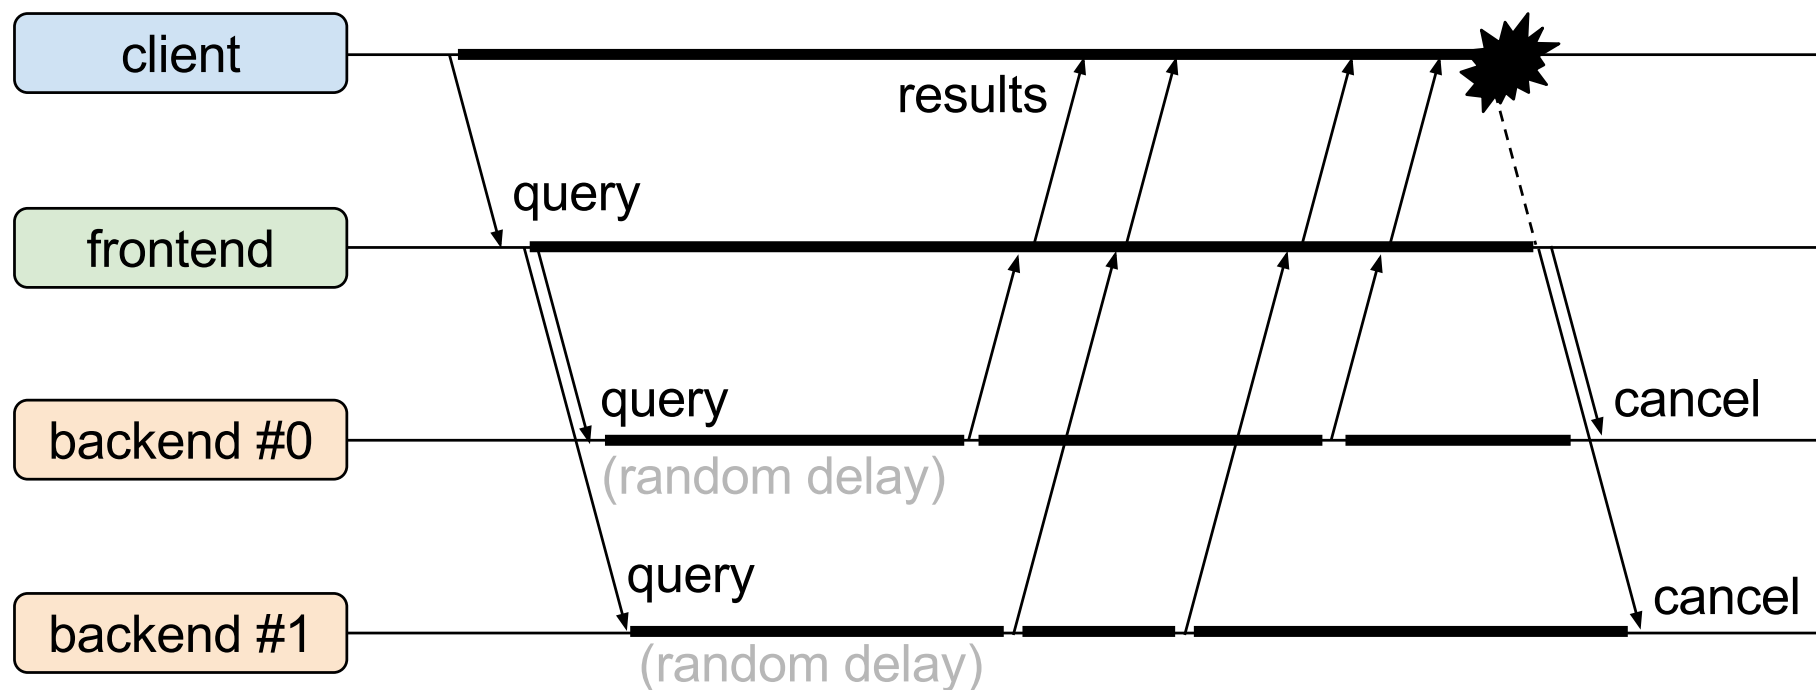
service Google {
    // Search returns a Google search result for the query.
    rpc Search(Request) returns (Result) {}

    // Watch returns a stream of Google search results for the query.
    rpc Watch(Request) returns (stream Result) {}
}

message Request {
    string query = 1;
}

message Result {
    string title = 1;
    string url = 2;
    string snippet = 3;
}
```


Frontend uruchamia Watch na obu serwerach backend



Demo client --mode=watch

- Debugowanie aktywnego strumienia
- Anulowanie żądania


Kod źródłowy - client



Kod źródłowy - client (watch)

```
public void watch(String query) {  
    final Request request = Request.newBuilder().setQuery(query).build();  
    final CountDownLatch latch = new CountDownLatch(1); // we expect only 1 result  
  
    StreamObserver<Result> stream = new StreamObserver<Result>() {  
        @Override  
        public void onNext(Result value) {  
            logger.info("Search result: " + value.getTitle());  
        }  
        @Override  
        public void onError(Throwable t) {  
            logger.severe(("Error while watching for results! " + t.getMessage()));  
            latch.countDown();  
        }  
        @Override  
        public void onCompleted() {  
            logger.info("Watch done!");  
            latch.countDown();  
        }  
    };  
  
    googleStub.watch(request, stream);  
    Uninterruptibles.awaitUninterruptibly(latch, 100, TimeUnit.SECONDS);  
}
```

Kod źródłowy - backend



Kod źródłowy - backend (watch)

```
public class Backend {  
    private int port = 36061;  
    private Server server;  
  
    private void start(int id) throws IOException {  
        server = ServerBuilder.forPort((port + id))  
            .addService(new GoogleImpl(id))  
            .build()  
            .start();  
        logger.info("Server started, listening on " + (port + id));  
    }  
  
    public static void main(String[] args) throws IOException, InterruptedException {  
        String id = "0";  
  
        if (args.length > 0) {  
            id = args[0];  
        }  
  
        final Backend server = new Backend();  
        server.start(Integer.parseInt(id));  
        server.blockUntilShutdown();  
    }  
}
```

Kod źródłowy - backend (watch)

```
class GoogleImpl extends GoogleGrpc.AbstractGoogle {
    private final int id;


    public GoogleImpl(int id) {
        this.id = id;
    }

    @Override
    public void watch(Request req, StreamObserver<Result> responseObserver) {
        int responseNo = 0;
        final ServerCallStreamObserver serverCall = (ServerCallStreamObserver) responseObserver;
        serverCall.setOnCancelHandler(() -> logger.warning("Request canceled!"));

        while (!serverCall.isCancelled()) {
            sleepUpToMiilis(1000);
            responseObserver.onNext(Result.newBuilder().setTitle(
                format("result %d for [%s] from backend %d", responseNo++, req.getQuery(), id)).build());
        }

        private void sleepUpToMiilis(int millis) {
            Try.ofFailable(() -> { Thread.sleep((generator.nextInt(millis / 10) + 1) * 10); }).get();
        }
    }
}
```

Kod źródłowy - frontend



Kod źródłowy - frontend (watch)

```
func (s *server) Watch(req *pb.Request, stream pb.Google_WatchServer) error {
    ctx := stream.Context()
    c := make(chan result)
    var wg sync.WaitGroup
    for _, b := range s.backends {
        wg.Add(1)
        go func(backend pb.GoogleClient) {
            defer wg.Done()
            watchBackend(ctx., backend, req, c)
        }(b)
    }
    go func() {
        wg.Wait()
        close(c)
    }()
    for res := range c {
        if res.err != nil {
            return res.err
        }
        if err := stream.Send(res.res); err != nil {
            return err
        }
    }
    return nil
}
```


Kod źródłowy - frontend (watchBackend)

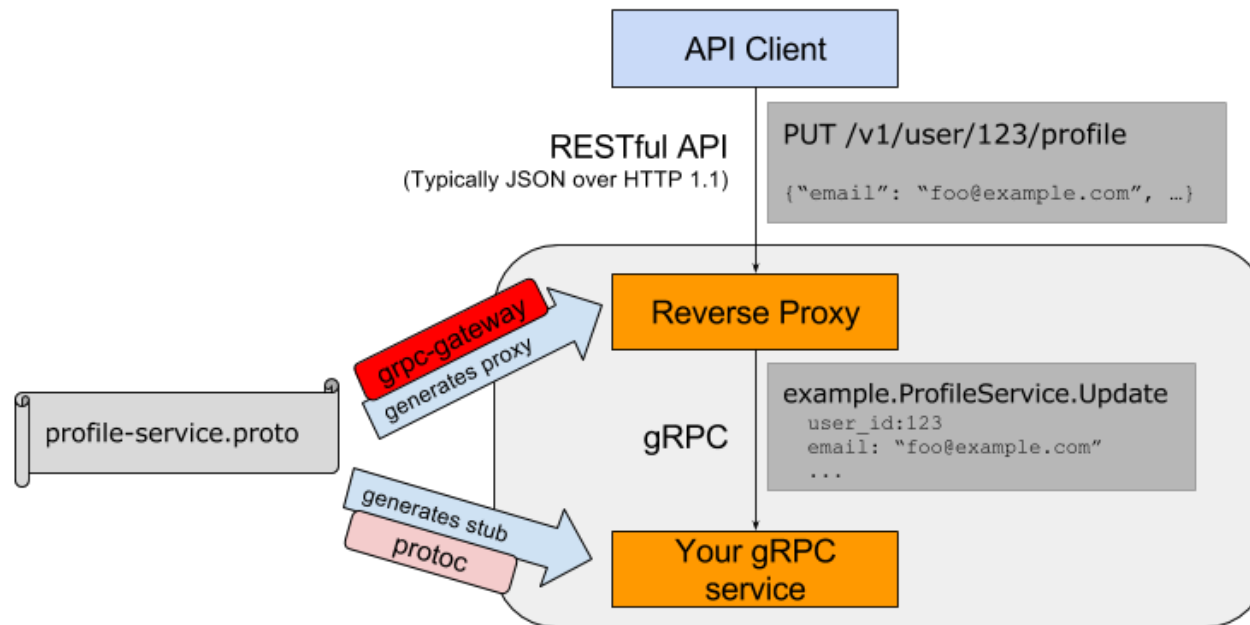
Watch kończy swoje zadanie przy pierwszym napotkanym błędzie; to anuluje żądanie za pomocą `ctx.Done` oraz kończy wywołanie metody `watchBackend`.

```
func watchBackend(ctx context.Context, backend pb.GoogleClient, req *pb.Request, c chan<- result)
    stream, err := backend.Watch(ctx, req)
    if err != nil {
        select {
            case c <- result{err: err}:
            case <-ctx.Done():
            }
        return
    }
    for {
        res, err := stream.Recv()
        select {
            case c <- result{res, err}:
                if err != nil { return }
            case <-ctx.Done():
                return
        }
    }
}
```

gRPC Gateway - github.com/gengo/grpc-gateway

grpc-gateway to plugin do protoc, który wczytuje definicje serwisów z plików *.proto i tworzy reverse-proxy, który tłumaczy RESTful JSON API na gRPC.

Pomaga tworzyć jedno API, które dostępne jest przy pomocy klienta gRPC oraz za pomocą RESTa jednocześnie.



gRPC podsumowanie

- Idealny do rozproszonych systemów (strumienie, anulowanie żądań, śledzenie)
- Szybki - Prototol Buffers
- Samodokumentujący - RPC
- HTTP/2
- Wsparcie dla wielu języków(**Java**, **Go**, **C**, **C++**, Node.js, Python, Ruby, Objective-C, PHP, i C#)
- Idealny dla mobilnych aplikacji - Java, Objective-C
- Load balancer
- SSL/TLS

Pytania?

Thank you

Mateusz Dymiński

[@m_dyminski](https://twitter.com/m_dyminski) (http://twitter.com/m_dyminski)

github.com/mateuszdyminski/grpc

